

# **Advanced Statistics Assignment 2: Sampling Distribution & Estimation**

Alli Ajagbe Olasunkanmi

## Question 1

1. Obtaining the sampling distribution of the three estimators for all three models for  $n = 20$

Defining the estimators:

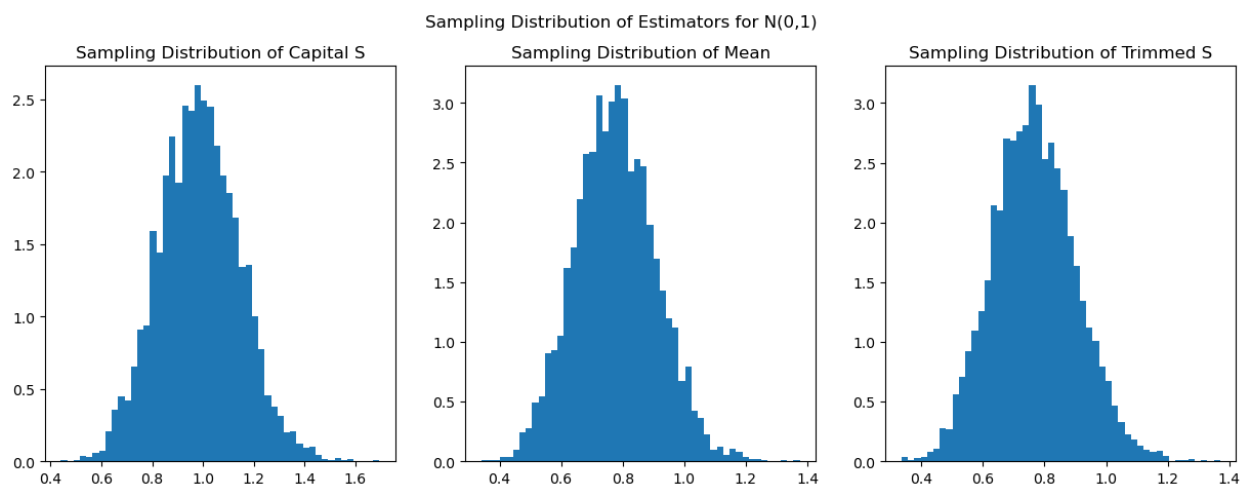
```
def myEstimators(n,x):  
    # sample standard deviation S  
    divisor = 1/(n - 1)  
    capital_S = np.sqrt(divisor * np.sum((x - np.mean(x))**2))  
  
    # mean deviation  
    mean = np.mean(abs(x - np.mean(x)))  
  
    # trimmed standard deviation  
    # discarding the smallest 0.1n/2 and largest 0.1n/2  
    new_x = np.sort(x)  
    new_x = new_x[int(0.1*n/2):int(n - 0.1*n/2)]  
    trimmed_S = np.sqrt(divisor * np.sum((new_x - np.mean(new_x))**2))  
  
    return capital_S, mean, trimmed_S
```

Defining the function to compute the estimators for  $N(0,1)$ :

```
n = 20  
  
def normal(n, iter=5000):  
    # creating lists to store the estimators  
    capital_Ss = []  
    means = []  
    trimmed_Ss = []  
  
    for i in range(iter):  
        # drawing a sample of size n from N(0,1)  
        x = np.random.normal(0,1,n)  
  
        # computing the estimators  
        capital_S, mean, trimmed_S = myEstimators(n,x)  
  
        capital_Ss.append(capital_S)  
        means.append(mean)  
        trimmed_Ss.append(trimmed_S)  
    return capital_Ss, means, trimmed_Ss  
  
capital_Ss, means, trimmed_Ss = normal(n)
```

## Plotting the values:

```
# plotting the sampling distribution of the estimators
fig, ax = plt.subplots(1,3,figsize=(15,5))
ax[0].hist(capital_Ss, bins=50, density=True)
ax[0].set_title('Sampling Distribution of Capital S')
ax[1].hist(means, bins=50, density=True)
ax[1].set_title('Sampling Distribution of Mean Deviations')
ax[2].hist(trimmed_Ss, bins=50, density=True)
ax[2].set_title('Sampling Distribution of Trimmed S')
plt.suptitle('Sampling Distribution of Estimators for N(0,1)')
plt.show()
```



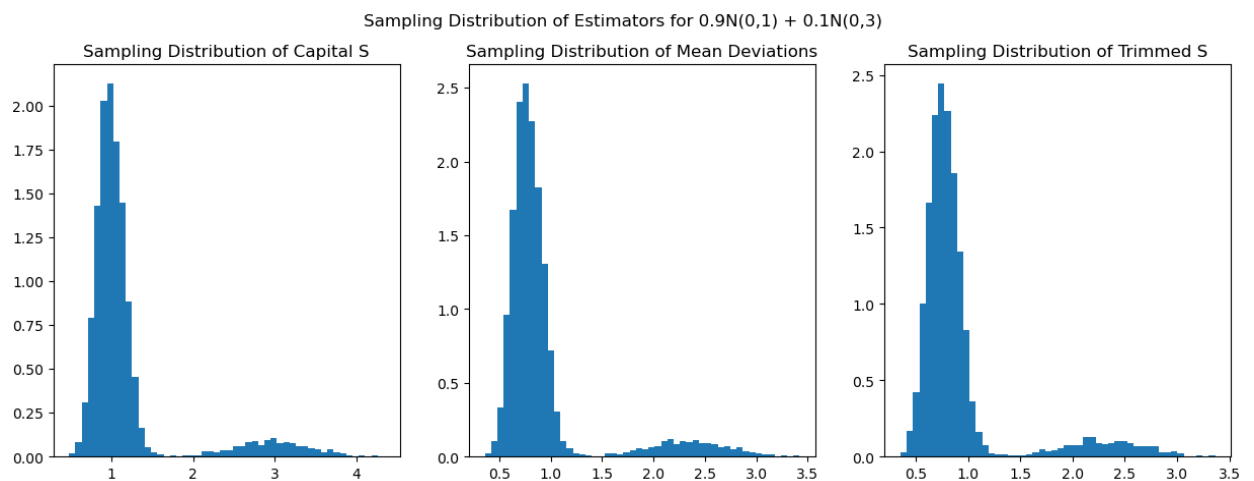
## Defining the function for Mixture of Two Normals: $0.9N(0,1) + 0.1N(0,1)$

```
def mixture_of_normals(var1, var2, n, iter=5000):
    capital_Ss = []
    means = []
    trimmed_Ss = []
    for i in range(iter):
        # selecting N(0,1) with probability 0.9 and N(0,3) with probability 0.1
        epsilon = 0.1
        if np.random.uniform() < epsilon:
            x = np.random.normal(0, var2, n)
        else:
            x = np.random.normal(0, var1, n)

        capital_S, mean, trimmed_S = myEstimators(n,x)
        capital_Ss.append(capital_S)
        means.append(mean)
        trimmed_Ss.append(trimmed_S)
    return capital_Ss, means, trimmed_Ss
```

## Plotting the values:

```
capital_Ss, means, trimmed_Ss = mixture_of_normals(1, 3, n)
fig, ax = plt.subplots(1,3,figsize=(15,5))
ax[0].hist(capital_Ss, bins=50, density=True)
ax[0].set_title('Sampling Distribution of Capital S')
ax[1].hist(means, bins=50, density=True)
ax[1].set_title('Sampling Distribution of Mean Deviations')
ax[2].hist(trimmed_Ss, bins=50, density=True)
ax[2].set_title('Sampling Distribution of Trimmed S')
plt.suptitle('Sampling Distribution of Estimators for 0.9N(0,1) + 0.1N(0,3)')
plt.show()
```



## Defining the function for the Mixture of a Normal and a Cauchy: 0.95N(0,1) + 0.05C(0,1)

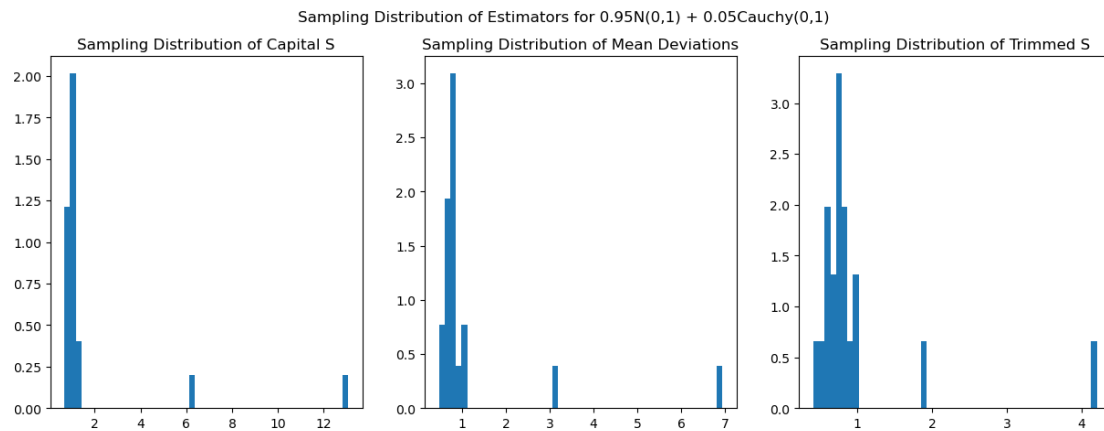
```
def normal_plus_cauchy(var1, iter=5000):
    capital_Ss = []
    means = []
    trimmed_Ss = []
    for i in range(iter):
        epsilon = 0.05
        if np.random.uniform() < epsilon:
            x = np.random.standard_cauchy(n)
        else:
            x = np.random.normal(0, var1, n)

        capital_S, mean, trimmed_S = myEstimators(n,x)
        capital_Ss.append(capital_S)
        means.append(mean)
        trimmed_Ss.append(trimmed_S)

    return capital_Ss, means, trimmed_Ss
```

## Plotting the values:

```
capital_Ss, means, trimmed_Ss = normal_plus_cauchy(1, n)
fig, ax = plt.subplots(1,3,figsize=(15,5))
ax[0].hist(capital_Ss, bins=50, density=True)
ax[0].set_title('Sampling Distribution of Capital S')
ax[1].hist(means, bins=50, density=True)
ax[1].set_title('Sampling Distribution of Mean Deviations')
ax[2].hist(trimmed_Ss, bins=50, density=True)
ax[2].set_title('Sampling Distribution of Trimmed S')
plt.suptitle('Sampling Distribution of Estimators for 0.95N(0,1) + 0.05Cauchy(0,1)')
plt.show()
```



## 2. Computing the Bias and MSE for all three Estimators

```
# computing the bias and mse of the estimators
def bias(estimator, true_value):
    return np.mean(estimator) - true_value

def mse(estimator, true_value):
    return np.var(estimator) + (np.mean(estimator) - true_value)**2
```

### *Normal Distribution:*

	Sample Standard Deviation	Mean Deviation	Trimmed Standard Deviation
<b>Bias</b>	-0.013679794574373627	-0.22215247287238782	-0.23259808527178905
<b>Mean Squared Error</b>	0.02620950081222507	0.067398138082861	0.07318563277916038

### *Mixture of Two Normals:*

	Sample Standard Deviation	Mean Deviation	Trimmed Standard Deviation
<b>Bias</b>	0.19075423279563175	-0.061843133382412985	-0.07415280178651373
<b>Mean Squared Error</b>	0.434760883780574	0.2568051550083415	0.25455566262543544

### *Normal + Cauchy:*

	Sample Standard Deviation	Mean Deviation	Trimmed Standard Deviation
<b>Bias</b>	-0.048951144364734134	-0.24569038948547195	-0.25254279256095613
<b>Mean Squared Error</b>	0.027067556147709962	0.0790549808160297	0.08414292164620482

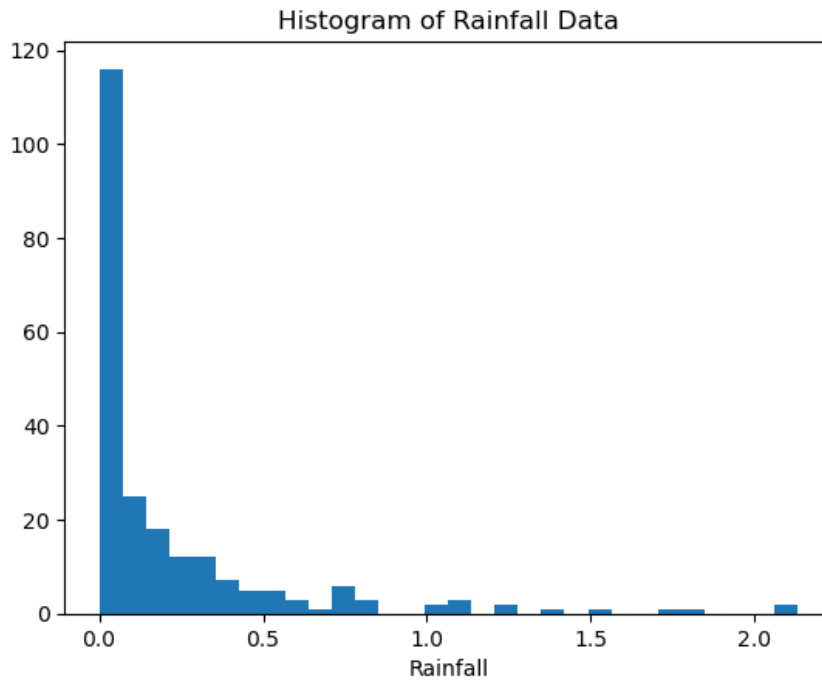
### **3. Observations**

Among the three estimators, the Sample Standard Deviation (S) tends to exhibit bias, especially in the presence of outliers and deviations from normality, resulting in higher variance and mean squared error (MSE). Mean Deviation (M) generally demonstrates lower bias compared to S across various data distributions and offers moderate resilience against outliers. Trimmed Standard Deviation typically features even lower bias than both S and M and is highly robust against outliers and non-normal data, making it a precise choice, often with the lowest variance and MSE among the three estimators. Hence, the choice and properties of estimators should be considered in relation to the data distribution available for inspection and the required estimation goals.

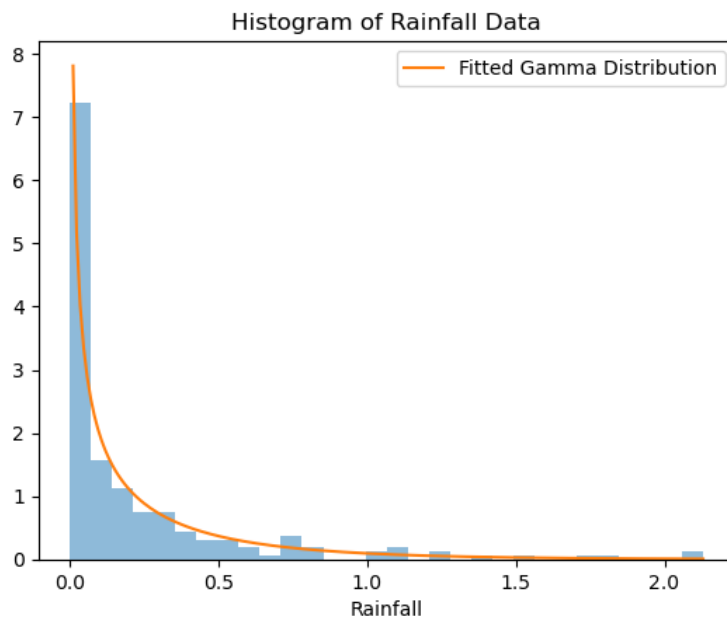
## Question 2

### 1. Fitting a Gamma distribution to the data:

The Data:



With fitted gamma plot:





a. **Estimating the parameters using the Method of Moments:**

```
# estimating the parameters using the method of moments
mean = np.mean(rainfall_data_flattened)
var = np.var(rainfall_data_flattened)
shape = mean**2/var
scale = var/mean
print("Shape:", shape)
print("Scale:", scale)
```

With the above, we get **0.3758789321633129** for *alpha* and **0.5966781904246276** for *beta*.

**Estimating the parameters using the Method of Maximum Likelihood with SciPy** because the inbuilt function uses MLE to fit the curve to the data:

```
# estimating the parameters using the method of maximum likelihood using scipy
alpha, loc, beta = stats.gamma.fit(rainfall_data_flattened, floc=0)
print("Alpha:", alpha)
print("Beta:", beta)
```

With the above, we get **0.43912250445350176** for *alpha* and **0.5107430359121928** for *beta*.

**Assessing the Goodness of Fit of the Gamma Distribution to the Data:**

This can be inspected visually by looking at the overlay of the curve on the actual data as plotted in the initial part of this question. We can proceed to visualise how well the gamma distribution fits the data by using a QQ Plot.

**Using the estimates from the Method of Moments:**

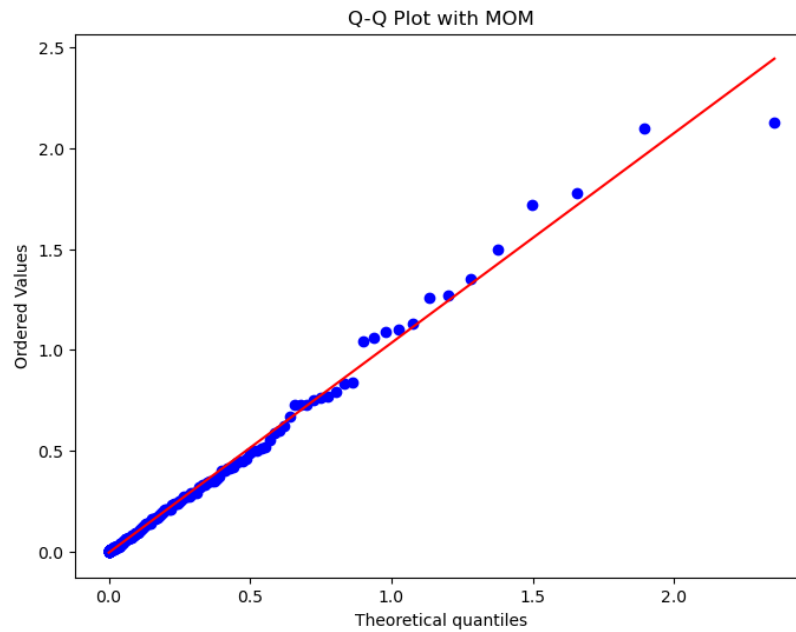
```
shape_value, scale_value = shape, scale
print(shape_value, scale_value)

# turning rainfall data to numpy array
rainfall_data = np.array(rainfall_data)

rainfall_data_flattened = rainfall_data.flatten()

# Create the Q-Q plot using stats.probplot
plt.figure(figsize=(8, 6))
res = stats.probplot(rainfall_data_flattened, dist=stats.gamma, sparams=(shape_value, 0, scale_value), plot=plt)
plt.title('Q-Q Plot with MOM')
plt.show()
```

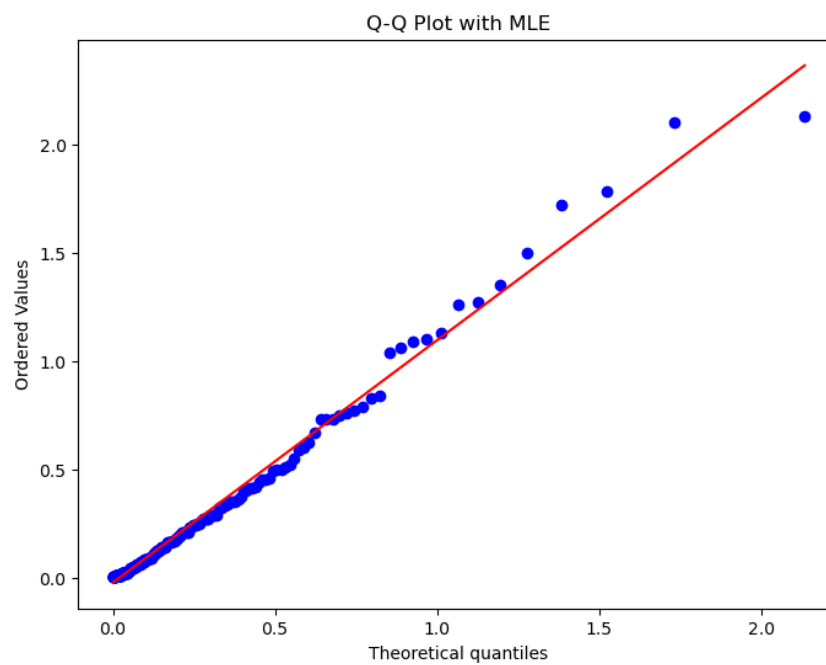
### Resulting QQ Plot:



### Using estimates from the Method of Maximum Likelihood:

```
shape_value, scale_value = alpha, beta
plt.figure(figsize=(8, 6))
res = stats.probplot(rainfall_data_flattened, dist=stats.gamma, sparams=(shape_value, 0, scale_value), plot=plt)
plt.title('Q-Q Plot with MLE')
plt.show()
```

### Resulting QQ Plot:




## 4. Parametric Bootstrap Method


### a. Using Maximum Likelihood Estimators

```
# finding the mle of the sample using scipy
alphas = []
betas = []
for i in range(1000):
    sample = stats.gamma.rvs(alpha, loc, beta, size=226)
    alpha_sample, loc_sample, beta_sample = stats.gamma.fit(sample, floc=0)
    alphas.append(alpha_sample)
    betas.append(beta_sample)

# plotting the histogram of the mle of the sample
sns.histplot(alphas, bins=50, color='red')
plt.title('Histogram of Alpha')
plt.xlabel('Alpha')
plt.show()

sns.histplot(betas, bins=50, color='green')
plt.title('Histogram of Beta')
plt.xlabel('Beta')
plt.show()
```



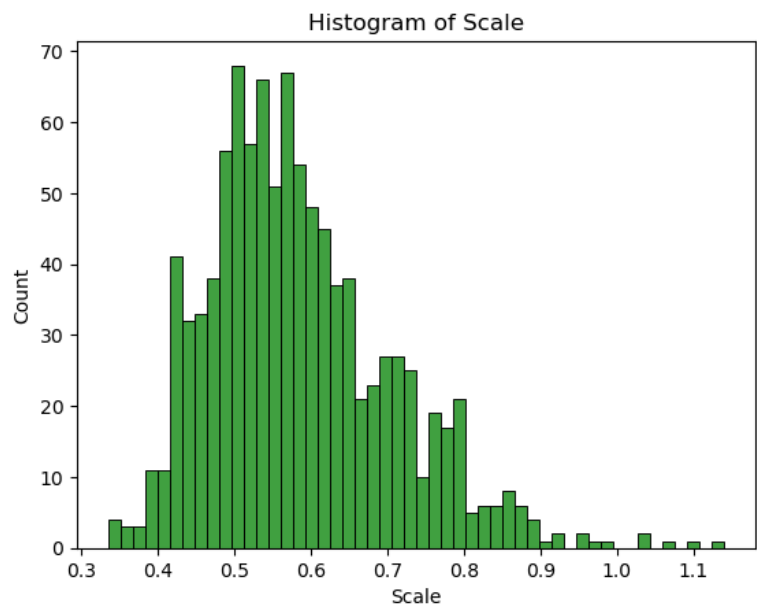
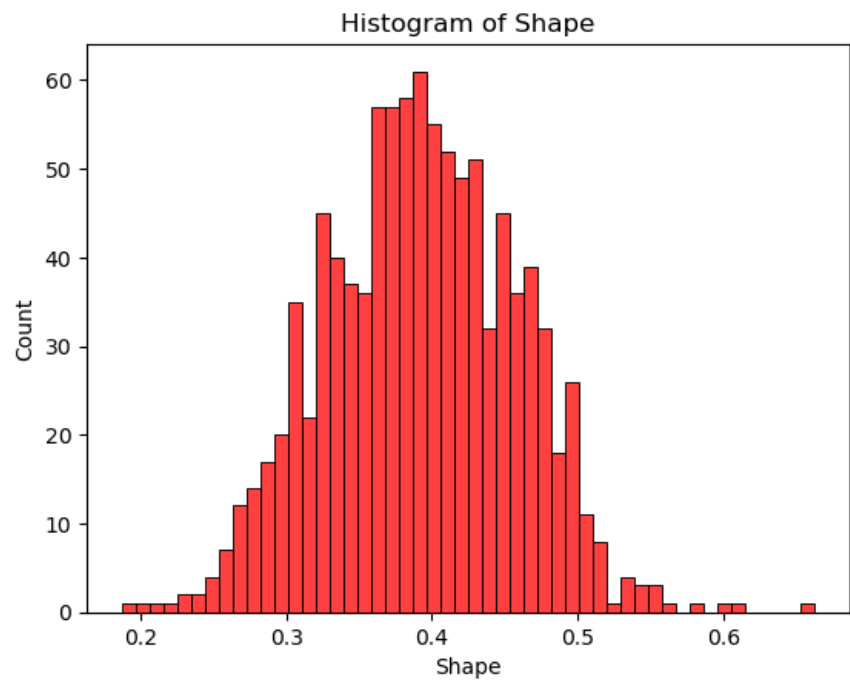


## b. Using Methods of Moment Estimators

```
# finding the mom estimators of the sample
shapes = []
scales = []
for i in range(1000):
    sample2 = stats.gamma.rvs(shape, loc, scale, size=226)
    mean_sample = np.mean(sample2)
    var_sample = np.var(sample2)
    shape_sample = mean_sample**2/var_sample
    scale_sample = var_sample/mean_sample
    shapes.append(shape_sample)
    scales.append(scale_sample)

# plotting the histogram of the mom estimators of the sample
sns.histplot(shapes, bins=50, color='red')
plt.title('Histogram of Shape')
plt.xlabel('Shape')
plt.show()

sns.histplot(scales, bins=50, color='green')
plt.title('Histogram of Scale')
plt.xlabel('Scale')
plt.show()
```



### c. Comparing Bias and MSEs

```
# comparing the mle and mom estimators using the bias and mse
print("Bias of Alpha MLE:", np.mean(alphas) - shape)
print("Bias of Beta MLE:", np.mean(betas) - scale)
print("Bias of Alpha MOM:", np.mean(shapes) - shape)
print("Bias of Beta MOM:", np.mean(scales) - scale)

# using var plus bias squared to calculate mse
print("MSE of Alpha MLE:", np.var(alphas) + (np.mean(alphas) - shape)**2)
print("MSE of Beta MLE:", np.var(betas) + (np.mean(betas) - scale)**2)
print("MSE of Alpha MOM:", np.var(shapes) + (np.mean(shapes) - shape)**2)
print("MSE of Beta MOM:", np.var(scales) + (np.mean(scales) - scale)**2)
```

	MLEs	MoMs
<b>Bias - Alpha</b>	0.0038493895265862 554	0.01775848273622765 4
<b>Bias - Beta</b>	-0.0052954500338473 49	-0.0109005179439667 37
<b>MSE - Alpha</b>	0.00108127932498524 7	0.00626227355463753 8
<b>MSE - Beta</b>	0.00340772993514698 54	0.014519833341281338

Based on the information in the table above, it is evident that Maximum Likelihood Estimation (MLE) performs better than the Method of Moments (MoM) in estimating the parameters. This conclusion is supported by two key observations:

**Bias Comparison:** The bias values for both Alpha and Beta are lower for MLE compared to MoM. This indicates that MLE provides more accurate parameter estimates, as lower bias values suggest estimates closer to the true values.

**Mean Squared Error (MSE) Comparison:** The MSE values for both Alpha and Beta are also lower for MLE compared to MoM. This suggests that MLE is more precise in estimating the parameters, as lower MSE values indicate estimates with smaller overall errors.

